


Living in a high
dimensional world

Hello there

- Michael Garcia Ortiz – michael.garcia-ortiz@city.ac.uk
 - New Lecturer
 - Former Research Scientist in Robotics company
 - Former consultant in ML (finance, startups)
- Interests: Robotics
Artificial General Intelligence
(I have UG – PG project ideas)



Goal: give you an overview of what you can do with data.



What you need for this module


- Understanding how things work:
Mathematics <https://mml-book.com>
- Making things work:
git, python, sklearn, matplotlib with **Jupyter Lab**
- Having a feeling about what might work:
Practice, trial and error, gaining insights and intuition



kaggle

Do the thing, and you shall have the power.

(insert thunder)



What we will see together

- Introduction: Living in a high dimensional world
 - The effect of dimensionality
 - Solving problems with data
 - Feature engineering
 - Projection and visualization
- Finding Intrinsic Structure in Data with Unsupervised Learning
- Predicting with Supervised Learning

What is high dimensionality?



High dimensional data is ubiquitous



High dimension is a double-edge sword



Transforming raw data



Projecting and visualizing your dataset

High dimensionality is everywhere

- Facebook: 350 M pictures every day (in 2013)
- Starbucks: 90 M transactions every week (in 2018)
- Universities: Course Signals at Purdue
- Big data in governments, Finance, Health ...



What is dimensionality?



DNA?



Human?



VGA: 10^6




EigenFace: 100




Dimensionality always **depends on your level of description**.
Often, you don't control the dimensionality of the data you obtain.

- Data you observe is the result of physical process:
 - It is the projection of this process on your sensor (image, sound, csv file)
- Example: Relevant objects, relevant object features for grasping
- Assumption: data can be manipulated, compressed, abstracted, vectorized. Made more explicit, easier to handle




Intrinsic dimension




Data dimensionality?

Examples from Kaggle



- Los Angeles Parking Citations : 19
 - Make, color
 - Location
- Chest X-Ray Images (Pneumonia) : ~ 1400x1000
- New York City Airbnb Open Data : 16
 - Location
 - Room type
 - Price
 - Name

Exploitable dimensionality?



OFTEN NEED TO **MANIPULATE DATA**:
BINNING, ONE-HOT ENCODING,
FEATURE EXTRACTION, COMPRESSION,
ABSTRACTION, REPRESENTATION...



TABULAR DATA: ACTUAL DIMENSION IS
PROBABLY HIGHER THAN THE
DIMENSION OF YOUR SPREADSHEET



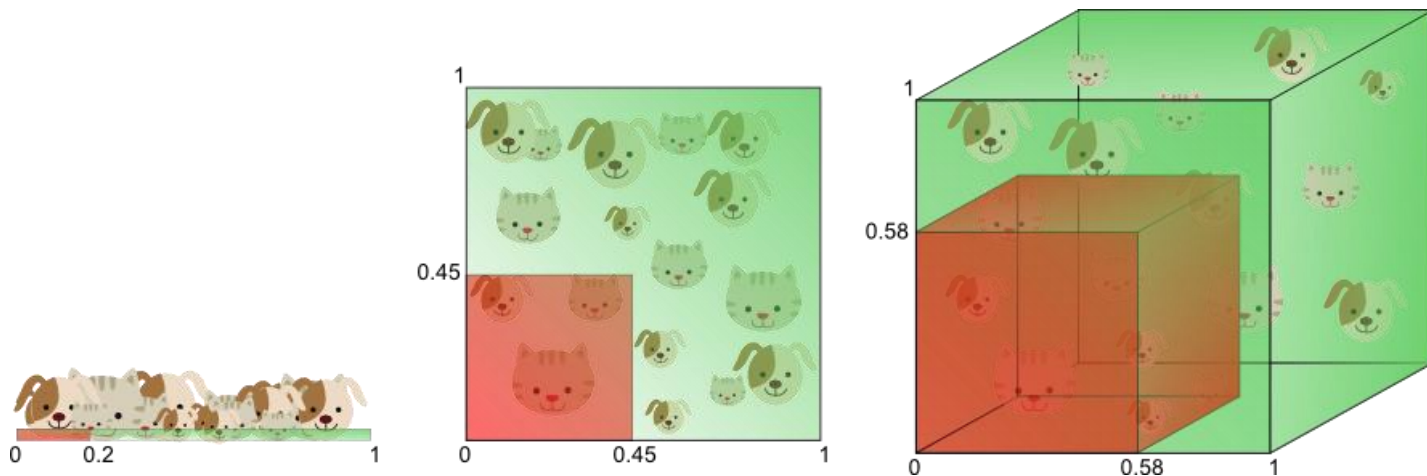
IMAGE DATA: ACTUAL DIMENSION IS
PROBABLY LOWER THAN THE
DIMENSION OF YOUR IMAGE

Why is high dimensionality a problem?

- **Visualization:** How do you communicate your findings?
Top management, non-technical people: We need very **communicative** visualization.
- **Interpretation:** Hypothesis about patterns in data you observe. How do you **confirm your insights**?
- **Difficult to process:** memory, compute. Multiplied by the number of models you want to try out, size of database, ...

The curse of dimensionality

- Higher dimensionality requires more data



High dimensionality is very beneficial

- More dimensions to express your problem

- Example: describe an animal!

- Number of legs
- Carnivorous?
- Predator?
- Size?
- Size of ears?



- It is far easier when the number of informative features is high.
- But, high dimensionality is not always equal to more information

How to describe a problem?

- It is the role of a Data Scientist to decide how to describe a problem
- Extract relevant information
- Visualize the data
- Find patterns to solve problems

Approaches of Datascience



Once we have data, what do we do with it?



Visualization



Compression



Find patterns



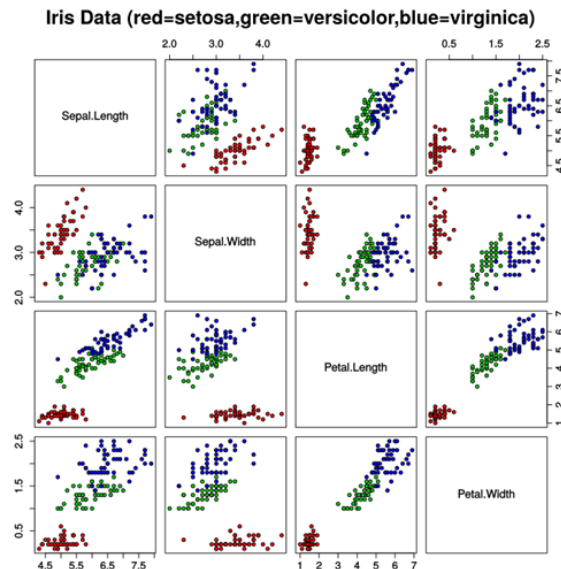

Solve problems

Dimensionality Reduction

- Storage and transfer
- Describing your problem correctly
- Abstraction to improve machine learning
- Other challenges: visualization, anonymity, transfer, ...

Visualization

- Why is it hard in high dimension?

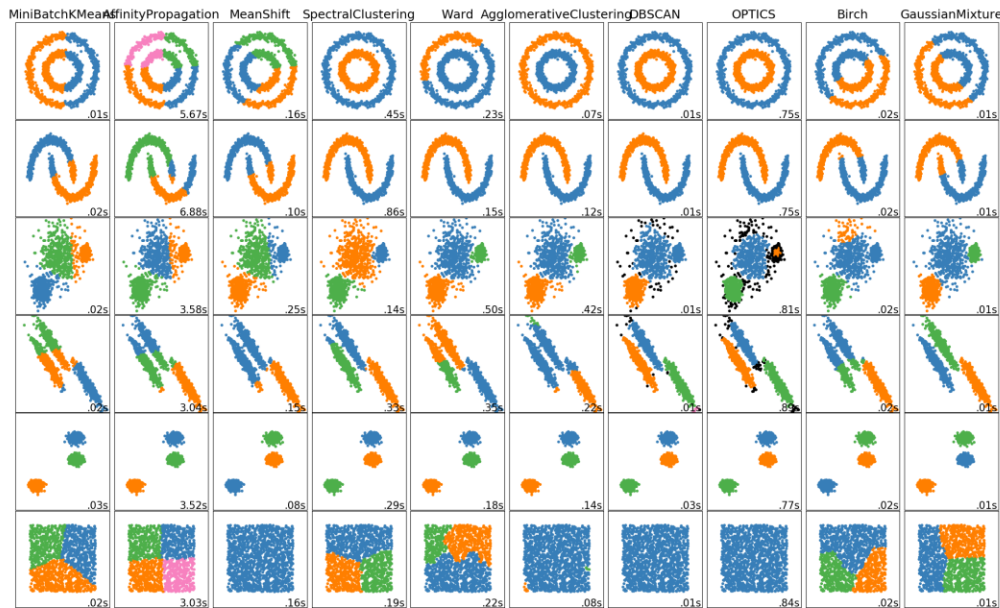


Find patterns

- Relations between dimensions of data
- Correlation, covariance, prediction
- Clusters, similarity / distance between data points
- Can data be used to predict things?


Unsupervised Learning: Clustering

- Grouping data points together by similarity
- What is similarity?
- How do you group?
- Assumptions?



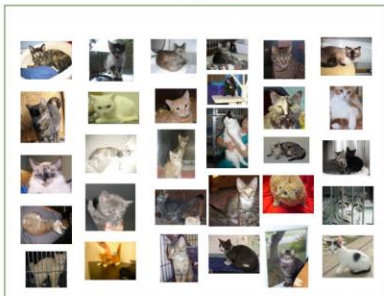
Supervised Learning: Regression

- Classic example: linear regression
- Input data, model, target
- But can be much more complex!

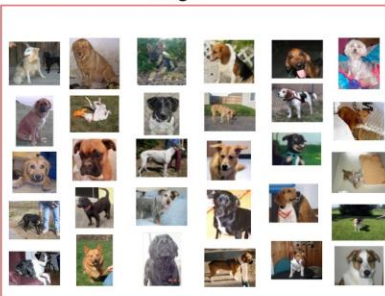


Supervised Learning: Classification


Cats



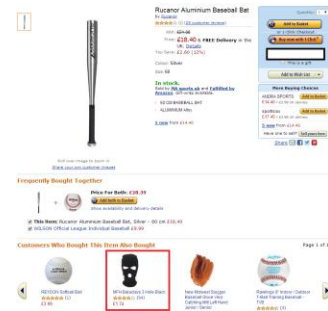

Dogs



Sample of cats & dogs images from Kaggle Dataset



Anomaly detection, recommender systems, ...



You need to represent your data with numbers

- Create features
- Input data is a curated list of features
- Algorithms will perform analysis based on the features you choose.

Feature engineering



Manage the dimensionality of the data



Bring your knowledge and intuitions to the problem:

Facilitates Machine Learning
Knowledge comes with bias





Dimensionality reduction through compression



Feature selection

Basic data preparation

- Continuous data: binning
- Categorical data: one-hot encoding
- Ordinal data: numerical
- Missing data: imputation
- Temporal data?
- Scaling, normalization, whitening...



Label Encoding

Food Name	Categorical #	Calories
Apple	1	95
Chicken	2	231
Broccoli	3	50

One Hot Encoding

Apple	Chicken	Broccoli	Calories
1	0	0	95
0	1	0	231
0	0	1	50



Why creating features?

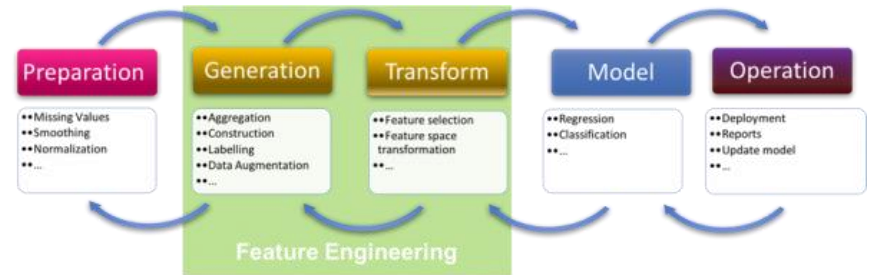
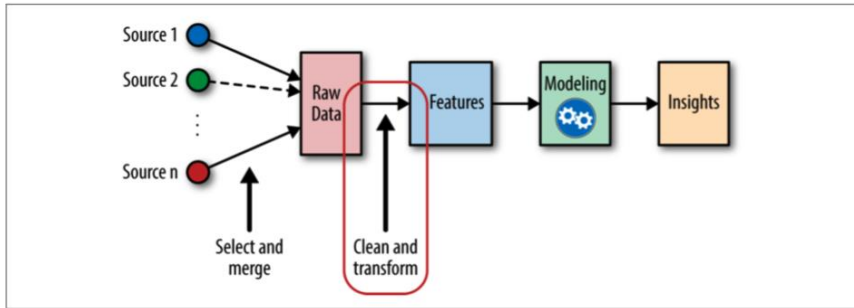


Importance of having indicators for change of trend, price drop, sudden price increase. Automatic alert? Autonomous trading?

Technical reasons



- Some ML algorithms depend on dimensionality of input data.
- Dimensionality too low: not expressive enough
- Dimensionality too high: can't compute anymore
- Benefit from more descriptive, easily separable features

Feature engineering








Bring domain knowledge to the table

- GPS coordinates on AirBnB listings: 2 real numbers
 - Why not distance to closest subway station? 1 number
 - Distance to landmarks? N numbers
- Sentiment analysis in political tweets
 - Political jargon can be very country-specific (Ex: US / UK)
- Predicting driver behavior:
 - distance to the next car
 - time to collision




Example: Time



- Timescales


- index

Example: Multimodal data




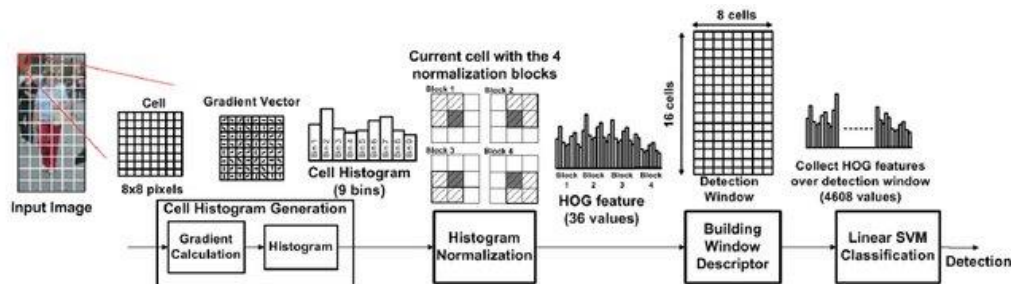
Student 1: 0.4 , -0.1

Student 2: 0.001, 0.5




San Francisco public transport

Example: Image descriptors



Example: Finance technical indicators

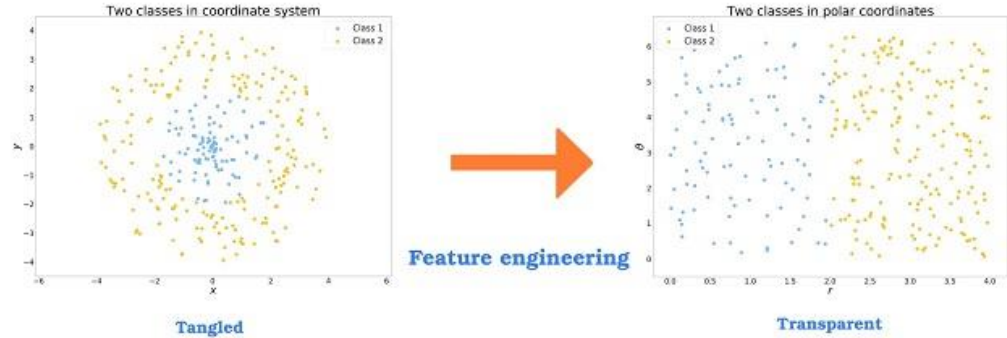


- Different kinds of traders: technical analysis, fundamental analysis...


- Technical analysis: look at signals, quantifiable indicators.

If these traders influence the market, then these signals become important to predict how the market will evolve.

Feature extraction





- Given current data representation, transform into other representation.
- More compact, more robust, more explicit input space
- Or just easier to handle for Machine Learning



Principal Component Analysis

- **Reduce** the dimensionality of a data set by finding a **new set of variables (i.e., principal components)**, smaller than the original set of variables
- Principal components (**PCs**) are **linear combinations** of existing variables
- Ordered by explained variance
- PCs are **uncorrelated**



PCA: Algorithm

- Center and normalize the data


- Compute covariance matrix

$$\Sigma_{ij} = \text{cov}(X_i, X_j) = E[(X_i - \mu_i)(X_j - \mu_j)]$$

- Calculate eigenvectors of covariance matrix -> PCs
- Project the data on the PCs
- How many dim do you keep?



PCA, details



PCA: example with Young People Survey



```
import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
```

```
young_ds = pd.read_csv('data/responses.csv')
```

```
young_ds_music = young_ds.iloc[:, :19].dropna()
young_ds_music.head()
```

		Slow songs or fast songs	Dance	Folk	Country	Classical music	Musical	Pop	Rock	Metal or Hardrock	Punk	Hip hop, Rap	Reggae, Ska	Swing, Jazz	Rock n roll	Alternative	Latino	Techno, Trance	Opera
0	5.0	3.0	2.0	1.0	2.0	2.0	1.0	5.0	5.0	1.0	1.0	1.0	1.0	1.0	3.0	1.0	1.0	1.0	1.0
1	4.0	4.0	2.0	1.0	1.0	1.0	2.0	3.0	5.0	4.0	4.0	1.0	3.0	1.0	4.0	4.0	2.0	1.0	1.0
2	5.0	5.0	2.0	2.0	3.0	4.0	5.0	3.0	5.0	3.0	4.0	1.0	4.0	3.0	5.0	5.0	5.0	1.0	3.0
3	5.0	3.0	2.0	1.0	1.0	1.0	1.0	2.0	2.0	1.0	4.0	2.0	2.0	1.0	2.0	5.0	1.0	2.0	1.0
4	5.0	3.0	4.0	3.0	2.0	4.0	3.0	5.0	3.0	1.0	2.0	5.0	3.0	2.0	1.0	2.0	4.0	2.0	2.0

- Music preferences (19 items)
- Movie preferences (12 items)
- Hobbies & interests (32 items)
- Phobias (10 items)
- Health habits (3 items)
- Personality traits, views on life, & opinions (57 items)
- Spending habits (7 items)
- Demographics (10 items)

```
data_music = young_ds_music.to_numpy()
pca = PCA(whiten = True)
X_pca = pca.fit_transform(data_music)
print( pca.explained_variance_ratio_)
```

```
[0.21462856 0.15159296 0.10932881 0.06389559 0.05871546 0.05128242
0.04683182 0.03940256 0.03736163 0.03154765 0.03065362 0.02501457
0.02351328 0.02224939 0.0209301 0.02061214 0.01940997 0.0187642
0.01426525]
```



```
young_ds_hobbies = young_ds.iloc[:, 19+12:19+12+32].dropna()
data_hobbies = young_ds_hobbies.to_numpy()
pca = PCA(whiten = True)
X_pca = pca.fit_transform(data_hobbies)
print( pca.explained_variance_ratio_)
```

```
[0.13544799 0.10479068 0.08063201 0.06946438 0.05273425 0.04286766
0.04094587 0.03457081 0.0330224 0.03203322 0.02858578 0.02790116
0.02637219 0.02556872 0.02287498 0.02230709 0.02090546 0.01967975
0.01890299 0.01870482 0.01676079 0.01558494 0.01492138 0.01457642
0.0131472 0.01233027 0.01150779 0.01002716 0.00894522 0.00866793
0.00825431 0.00696436]
```

```
young_ds_movies = young_ds.iloc[:, 19:19+12].dropna()
data_movies = young_ds_movies.to_numpy()
pca = PCA(whiten = True)
X_pca = pca.fit_transform(data_movies)
print( pca.explained_variance_ratio_)
```

```
[0.2410899 0.16954119 0.13430844 0.08476939 0.07149352 0.06382178
0.0593179 0.05306483 0.04439523 0.0290464 0.02663497 0.02251645]
```

Linear Discriminant Analysis



- LDA reduces dimensionality while preserving as much of the **class discriminatory information** as possible
- Suitable when you have class labels and want to find a projection that will help you to separate them
- Tries to **maximise** the difference between **group means** while **minimizing** the **within-class variance**
- **Maximize:**
$$J(w) = \frac{|\tilde{\mu}_1 - \tilde{\mu}_2|^2}{\tilde{s}_1^2 + \tilde{s}_2^2}$$
- Can be generalized to high dimension and multi-class




LDA: example

32	History
33	Psychology
34	Politics
35	Mathematics
36	Physics
37	Internet
38	PC Software, Hardware
39	Economy, Management
40	Biology
41	Chemistry
42	Poetry reading
43	Geography
44	Foreign languages
45	Medicine
46	Law
47	Cars
48	Art
49	Religion
50	Outdoor activities
51	Dancing

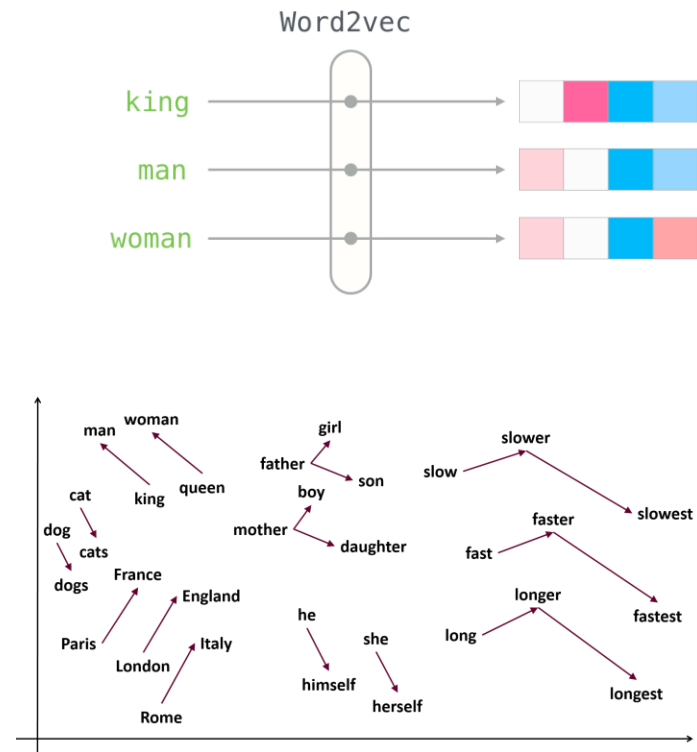
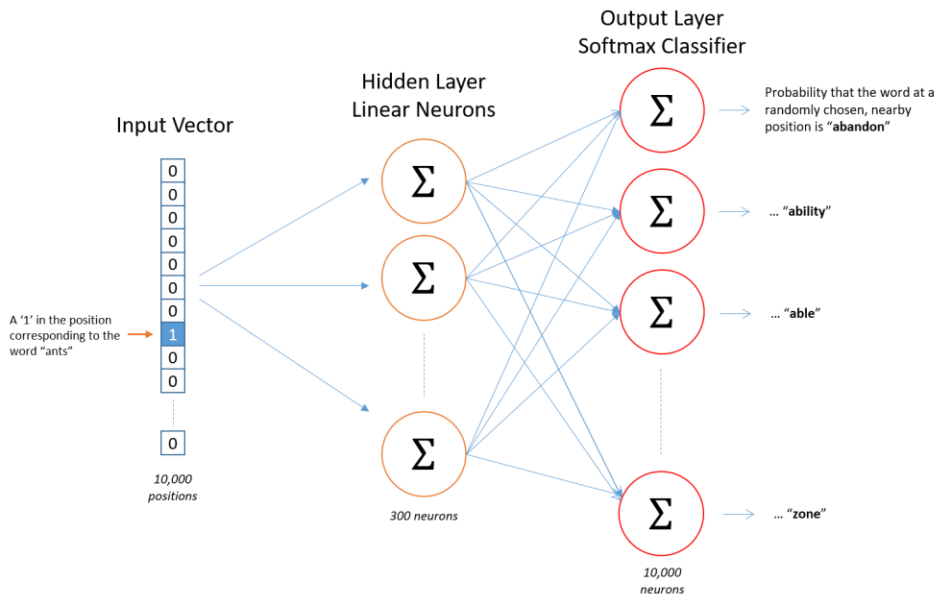
```
pca = PCA(whiten = True)
X_pca = pca.fit_transform(data_hobbies)

young_ds_drinking = young_ds.iloc[:, 74]
y_target, target_names = young_ds_drinking.factorize()



lda = LinearDiscriminantAnalysis(n_components=2)
lda_projection = lda.fit(data_hobbies, y_target).transform(data_hobbies)
```



Word2Vec



MFCCs




Feature selection

- Choose best features
- Limit number to prevent computational issues
- Measure quality of feature
- Filter vs wrapper vs embedded
- Other topics: Feature expansion, Automatic feature generation


Filter:



Wrapper:







Embedded:



From: https://en.wikipedia.org/wiki/Feature_selection

Deep Learning?

- Leverages huge quantity of data to **learn features**



- But you won't always have big data problems!


- Dataset size and dimensionality
- Memory, compute



- Finance: need insurance, cover risks
- Medical: help doctors, not replace them
- Self-driving cars (don't worry, my algorithm was trained in Paris)
- Being able to know why something has gone terribly wrong...



Sometimes, you need explainable features



Projection for visualization



Get intuition from data




Goal: project to 2D or 3D



Focus on non-linear models

Example 1: Multi-dimensional Scaling

	London	Berlin	Oslo	Moscow	Paris	Rome	Beijing	Istanbul	Gibraltar	Reykjavik
London	—									
Berlin	570	—								
Oslo	710	520	—							
Moscow	1550	1000	1020	—						
Paris	210	540	830	1540	—					
Rome	890	730	1240	1470	680	—				
Beijing	5050	4570	4360	3600	5100	5050	—			
Istanbul	1550	1080	1520	1090	1040	850	4380	—		
Gibraltar	1090	1450	1790	2410	960	1030	6010	1870	—	
Reykjavik	1170	1480	1080	2060	1380	2040	4900	2560	2050	—



MDS: algorithm




- Tries to preserve high-dimensional distances in the lower dimensional projection
- For instance, metricMDS (an implementation) tries to minimize:

$$\min_Y \sum_{i=1}^n \sum_{j=1}^n (d_{ij}^{(X)} - d_{ij}^{(Y)})^2$$

The diagram illustrates the minimization of the squared difference between two distance matrices. The equation $\min_Y \sum_{i=1}^n \sum_{j=1}^n (d_{ij}^{(X)} - d_{ij}^{(Y)})^2$ is shown. Two arrows originate from the terms $d_{ij}^{(X)}$ and $d_{ij}^{(Y)}$ in the equation. The arrow from $d_{ij}^{(X)}$ points to a grey box labeled "Original distances in high-dim.". The arrow from $d_{ij}^{(Y)}$ points to a grey box labeled "Distances in reduced dimensionality".

- Suitable for visualisation
- Depends on how distance is defined
- Curse of dimensionality (??)
- Meaning in **proximity** but not so much the axes.

MDS: example



Example 2: t-SNE

- Applicable to very large datasets
- Very effective to learn local structures
- Well-suited for 2D or 3D visualization

General idea: Match probability distribution of distances in high dimension and low dimension





Figure 25. Projection on the first two principal components of the latent distribution for the **Multi-dSprites** dataset. Each dot represents one object latent and is colored according to the corresponding ground truth factor.




Figure 26. Projection on the first two principal components of the latent distribution for the **Tetris** dataset. Each dot represents one object latent and is colored according to the corresponding ground truth factor.




Figure 27. t-SNE of the latent distribution for the **CLEVR6** dataset. Each dot represents one object latent and is colored according to the corresponding ground truth factor






Figure 28. t-SNE of the latent distribution for the **Multi-dSprites** dataset. Each dot represents one object latent and is colored according to the corresponding ground truth factor


Example 3: IsoMap




LLE projection




PCA projection




IsoMap projection





- MDS + geodesic distances
- Determine graph of neighbors
- Measure geodesic distances between nodes
- Compute dimensionality reduction on distances (MDS)



Example 4: SOMs

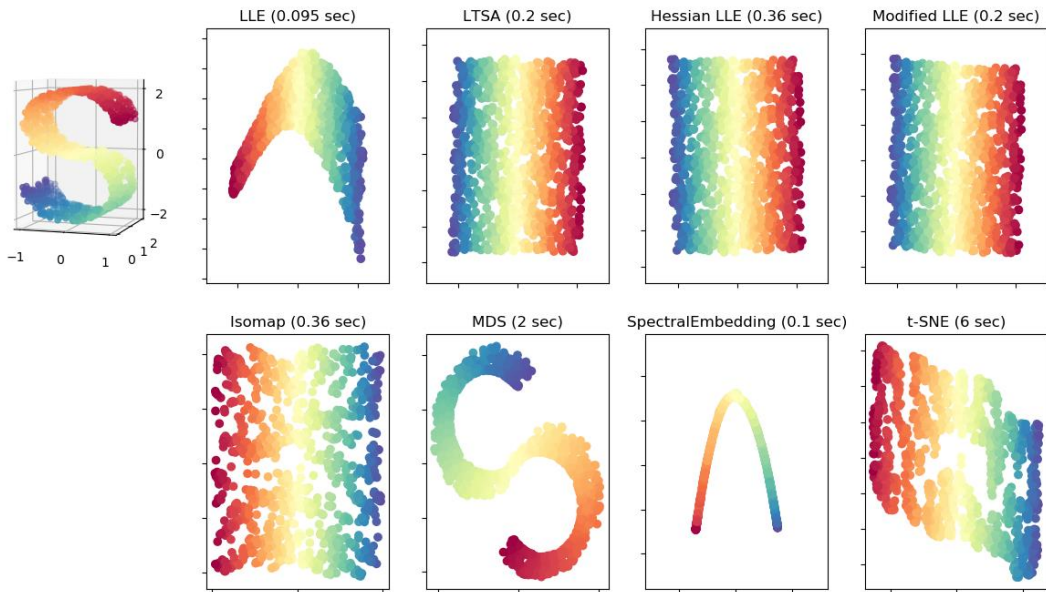


- Projection to 2D or 3D map using Neural Networks
- Iterative competitive learning:
 - Pick new data point
 - Find Best Matching Unit
 - Update BMU and neighbours toward data point



Other examples



Manifold Learning with 1000 points, 10 neighbors



- What are the assumptions on the shape of the manifold?
- Can it deal with the dimensionality of your dataset?

Feature learning and Big data ?



- Use huge quantity of data to learn features:
 - Plenty of labeled data: end-to-end
 - Plenty of unlabeled data: auto-encoder, generative models
 - Transfer: pre-trained network as feature extractor




- Difficult to visualize / interpret -> explainability
Open questions in AI:

- Human brain is (still) not explainable, but we still rely on it
- Should we rely on algorithms which perform better, but we don't know why?
- Who needs explainability? The user ? The engineer?


Questions?



Lateral Displacement



Longitudinal Displacement



Orientation

